
Rules:

Duration: 3 hours

Hint: Look at the point spread - attack with points in mind

Permitted: Open notes/reference material (including online)

<i>Problem</i>	<i>Score</i>	<i>Total</i>
<i>1</i>		<i>5</i>
<i>2</i>		<i>5</i>
<i>3</i>		<i>5</i>
<i>4</i>		<i>5</i>
<i>5</i>		<i>5</i>
<i>6</i>		<i>5</i>
<i>7</i>		<i>20</i>
<i>8</i>		<i>10</i>
		<i>50</i>

NOTE: *READ and SIGN the next page*
Make sure you have all 10 pages of the exam

- University's Academic Integrity Policy is in effect during this exam.
 - <http://academicintegrity.depaul.edu/AcademicIntegrityPolicy.pdf>

3. Violations of Academic Integrity

Violations of academic integrity include, but are not limited to, the following categories:

3.1. Cheating

Any action that violates University norms or instructor guidelines for the preparation and submission of assignments. This includes, but is not limited to:

- *Copying from another student.*
- *Offering, accepting, or otherwise obtaining or facilitating unauthorized assistance from or for another student.*
- *Having someone take an exam or complete an assignment in one's place.*
- *Unauthorized accessing of exam materials.*
 - *Accessing, using or possessing unauthorized materials during exams or quizzes.*

Please sign that you understand rules and university guidelines:

Name: Joey Domino

Date: 11/20/2022

Num of total pages: 10

Problem 1: (10pts) Deadlocks

[True/False] **Which of the following are guidelines for avoiding deadlocks?**

- a) Create a lock hierarchy so that locks protecting low-level data are prioritized over locks protecting high-level data.
- b) Use `std::recursive_mutex` to allow multiple locks using the same mutex. Using a recursive mutex resolves the undefined behavior that occurs when attempting to acquire a lock on a `std::mutex` which has already been used for a lock. When using a recursive mutex, the user just needs to ensure all locks on the mutex are released before the thread terminates.
- c) While already in possession of a lock, avoid calling user-supplied code since the user-supplied code's actions are unknown and may attempt to acquire a lock.
- d) Avoid nested locks, which can be avoided by not acquiring a lock if you already possess one, and if multiple locks are required, do it in a single action using `std::lock`.
- e) Call `lock` and `unlock` directly on the mutex, and avoid using `std::lock` and `std::lock_guard` except for the most trivial code.
- f) Always lock two mutexes in the same order (such as A before B), and when the mutexes are protecting a separate instance of the same class, pass both mutexes into the constructor of the lock and use `std::lock_guard` with the `std::adopt_lock` argument.

Answer: fill in **TRUE/FALSE** below (Are guidelines – True, Are NOT guidelines - False)

A) **False**

B) **False**

C) **True**

D) **False**

E) **False**

F) **True**

Problem 2: (5pts) Condition_Variable

[True/False] **Which statements are correct concerning wait() and notify()?**

- a) The lock used by wait() needs to be locked before wait() is called because it will be automatically unlocked by wait().
- b) A thread blocked by wait_for() will only be unblocked when the condition variable is notified or the relative timeout duration expires.
- c) All threads that call wait(std::unique_lock& lock) on the same condition variable need to acquire the lock on the same mutex.
- d) When a thread calls wait(std::unique_lock& lock, Predicate stop_waiting), the thread will be at least blocked once.
- e) When using the predicate version of wait_for(), it only returns false if the predicate function still evaluates to false after the timeout expired.
- f) If there are more than one threads waiting for the condition when notify_one() is called, it is unspecified which of the threads will be unblocked.
- g) notify_one() can unblock a thread that started waiting just after the call to notify_one() was made.
- h) The notifying thread needs to hold the lock on the same mutex as the one held by the waiting thread(s).

Answer: fill in as **TRUE/FALSE** below: (correct – True, incorrect – False)

- A) **True**
- B) **True**
- C) **False**
- D) **True**
- E) **False**
- F) **True**
- G) **False**
- H) **True**

Problem 3: (5pts) Non-Safe code

[short answer] **What is the non-thread safe aspect of the method “GetNextCount”?**

```
#include <mutex>

class DoStuff
{
public:
    int &GetNextCount()
    {
        std::lock_guard<std::mutex> lk(mtx);
        count++;
        return count;
    }

private:
    int count;
    std::mutex mtx;
};
```

Answer: (short answer)

The non-thread safe aspect of GetNextCount is the lack of unlocking the lock. I would either put the section that uses the lock inside it's own scope or make sure to unlock the mutex at the end.

Problem 4: (5pts) Future/Promise

[True/False] **Which of the following are appropriate use of a promise and future?**

- a) Thread A gets an unknown number of orders from a client, thread B waits until A is told it has received all orders and sends its orders to B.
- b) Thread A stores an integer that counts upwards one at a time at an arbitrary rate, and every time this integer reaches a multiple of 1000, it sends a signal to thread B, causing it to wake up.
- c) Thread A and thread B share an integer that may only be modified in one thread at a time. This value may be changed multiple times.
- d) Thread A and thread B run side by side simultaneously. When thread A reaches a certain result, both threads must end.
- e) Thread A stores an integer that counts up an arbitrary amount every time it loops. When this integer lands on a clean multiple of 10 for the first time, thread B will begin.

- f) Thread A and thread B both end with a result that the main thread must read. Main only needs to be able to read these values once thread A and B end.

Answer: fill in as **TRUE/FALSE** below: (*appropriate – True, inappropriate – False*)

A) *True*

B) *True*

C) *False*

D) *False*

E) *True*

F) *True*

Problem 5: (5pts) Join/Detach()

[True/False] **Which statements are correct concerning join() and detach()?**

- a) You do not need to call join() or detach() to safely exit a thread.
- b) The thread of execution will wait for a thread object to finish when calling join().
- c) The thread of execution will wait for a thread object to finish when calling detach().
- d) Detach() allows the thread to operate independently from the thread that created it.
- e) A thread is still joinable when calling detach().
- f) One should check that a thread is joinable before calling detach().

Answer: fill in as **TRUE/FALSE** below: (*correct – True, incorrect – False*)

A) False

B) True

C) False

D) True

E) False

F) True

Problem 6: (5pts) Launching a Thread

[short answer] Which of the following **cannot** launch the thread correctly?

There is a class with a function call operator:

```
class background_task
{
public:
    void operator>() const
    {
        do_something();
        do_something_else();
    }
};
```

Scenarios:

- a) `background_task f;`
`std::thread my_thread(f);`
- b) `std::thread my_thread(background_task());`
- c) `std::thread my_thread{background_task()};`
- d) `std::thread my_thread((background_task()));`
- e) `std::thread my_thread([]{`
 `do_something();`
 `do_something_else();`
 `});`

Answer: (list all Letters(a-e) that **cannot** launch the thread **with** [1-3 sentence justification](#))

B C D

Calling `background_task()` as such won't execute the `()` operator. This is calling the default constructor, which in this case won't do anything. To make this work, you would first need to create the object as A does, then call the object to run the operator.

Problem 7: (20 pts) Thought Experiment

Write an application to read 3 different binary files.

- a.bin
- b.bin
- c.bin

Concatenate these files into one final file, out.bin

At windows DOS command prompt you would type:

`copy /b a.bin+b.bin+c.bin out.bin`

Assume:

- these files are in different directories, requiring seek times and the files are fragmented
 - so it takes a long time to load from file to memory
- that fread(...) can load a file from disc to memory
 - It can access the files and stream the data into memory in parallel
 - Some magical DMA is installed that allows this to happen
 - this is a theoretical problem

Instead of using DOS command you write a method:

```
void Concatenate( char *outFileName, char *AFileName,
                 char *BFileName, char *CFileName );
```

for the above example:

```
Concatenate( "out.bin", "a.bin", "b.bin", "c.bin");
```

How would you write this application for concurrency?

[*pseudo code / sketch* – sample code doesn't need to compile]

```
Concatenate(char * outFile, char* AFile, char* BFile, char* CFile)
```

```
{
    SharedData shareData;

    //All follow same setup, but find the different directories
    Process A(AFile, outFile, shareData);
    Process B(BFile, outFile, shareData);
    Process C(CFile, outFile, shareData);

    //These set this->thread = std::thread(std::ref(*this));
    //starts the threads.
    //Destructor calls join on the thread.
    tA.Launch();
    tB.Launch();
    tC.Launch();
}
```

```
Process::Process(char * _inFile, char* _outFile, ShareData _sd)
    : inFile(_inFile), outFile(_outFile), sd(_sd)
{
}

Void Process::operator>()()
{
    //Lock the thread with shared mutex so we don't deadlock
    Std::unique_lock<std::mutex> lock (r.mtx);

    While(inFile != empty)
    {
        //populate outfile with the data from inFile
        if(filename == Apath)
        {
            Sd.AFile.parse(inFile);
        }
        if(filename == Bpath)
        {
            Sd.BFile.parse(inFile);
        }

        if(filename == Cpath)
        {
            Sd.CFile.parse(inFile);
        }
    }

    //Not perfect, but to get the idea. Check for all threads to be done
    //U
    If(inFile.name == A)
    {
        Sd.status == ShareData::Status::A_DONE;
    }
    If(inFile.name == B)
    {
        Sd.status == ShareData::Status::B_DONE;
    }

    If(inFile.name == C)
    {
        Sd.status == ShareData::Status::C_DONE;
    }
}
```

```
    }

    If(all done)
    {
        Sd.status == ShareData::Status::ALL_DONE;
    }

    //We finished. Fill the output file with the processed data
    If(sd.status == ShareData::Status::ALL_DONE
    {
        //In order preferably.
        outFile.concat(sd.AFile);
        outFile.concat(sd.BFile);
        outFile.concat(sd.CFile);
    }

    //we finished
    Lock.unlock()
}
```

```
SharedData
{
Public:
    Enum Status
    {
        A_DONE,
        B_DONE,
        C_DONE,
        ALL_DONE
    }

    //Big4
    Constructor
    Copy constructor
    Assignment operator
    Destructor
    //data
    Std::mutex mtx;
    Std::condition_variable cv;
    Status status;

    Char* AFile;
    Char* BFile;
    Char* CFile;
}
```


Problem 8: (10 pts) Reflection

You came to class with some Multithreaded views(beliefs), describe how those ideas were reinforced by our multithreading assignments (Jetsons and Maze) or changed.

I believed that multithreading was a relatively simple process. "It's just saying 'open this thread and do some stuff, then close the thread.'" He thought naively. Turns out I was sort of right.

Working with jetsons and the maze showed me how working with multiple threads is a bit like working with while loops within separate classes. I'm not 100% sure if a thread can exist without some kind of while loop forcing it to run forever. With this thinking, the only difference between a normal loop and a thread is that the threads run as their own process until rejoined or killed. In my mind, that's the same as a class communicating with another while they both have infinite while loops running. It takes some type of flag to end the loop either way. The threads just require mutex locks to not collide with eachother when dealing with shared functions.

It was the different mutexes that really screwed up my thinking. It was fascinating to read about the different lock types, but that's also where my head started to spin. In the end, I stuck with a unique lock for every process I built out. It worked fine, but I would have liked to have more time messing with different lock types to get better understanding of them.

I believe my thinking before this class was mostly correct, but lacking in real understanding of the details. After taking this class, I think I understand multithreading enough now that I could make my own projects with multithreading without much hassle. I just won't be too fast. I won't be Googling "how to multithread." Instead, I'll be searching for reminders like what the different lock types are and their use cases. I think that's a huge improvement. I'm excited to work with these systems more. I love seeing systems work fast!

